



ΕΘΝΙΚΟ ΚΑΙ ΚΑΠΟΔΙΣΤΡΙΑΚΟ ΠΑΝΕΠΙΣΤΗΜΙΟ ΑΘΗΝΩΝ

**ΣΧΟΛΗ ΘΕΤΙΚΩΝ ΕΠΙΣΤΗΜΩΝ
ΤΜΗΜΑ ΠΛΗΡΟΦΟΡΙΚΗΣ ΚΑΙ ΤΗΛΕΠΙΚΟΙΝΩΝΙΩΝ**

ΠΤΥΧΙΑΚΗ ΕΡΓΑΣΙΑ

Doop-Soot: Parallel Fact Generation

Mouris Dimitris

Επιβλέπων: Smaragdakis Yannis, Associate Professor NKUA

ΑΘΗΝΑ

ΣΕΠΤΕΜΒΡΗΣ 2016

ΠΤΥΧΙΑΚΗ ΕΡΓΑΣΙΑ

Doop-Soot: Parallel Fact Generation

Mouris Dimitris

A.M.: 1115201200114

ΕΠΙΒΛΕΠΩΝ: Smaragdakis Yannis, Associate Professor NKUA

ΠΕΡΙΛΗΨΗ

Παραλληλοποίηση του Fact Generation του Doop. Το Doop χρησιμοποιείται για μπλαμπλαμπλ

ΘΕΜΑΤΙΚΗ ΠΕΡΙΟΧΗ: Τεκμηρίωση

ΛΕΞΕΙΣ ΚΛΕΙΔΙΑ: static program analysis, doop: fact generation, soot, πτυχιακές εργασίες, τμήμα πληροφορικής και τηλεπικοινωνιών

Πανεπιστήμιο Αθηνών

ABSTRACT

Παραλληλοποίηση του Fact Generation του Doop. Το Doop χρησιμοποιείται για μπλαμπλαμπλαμπλ

SUBJECT AREA: Documentation

KEYWORDS: static program analysis, doop: fact generation, soot, undergraduate thesis, dept. of informatics

University of Athens

Αφιέρωση σε κάποιους.

ACKNOWLEDGEMENTS

blahblahblahblah

blahblahblahblah

blahblahblahblah

ΠΕΡΙΕΧΟΜΕΝΑ

PROLOGUE	11
1. INTRODUCTION	12
2. DOOP	13
2.1 Fact Generation	13
2.2 Doop Time Examples	13
3. SOOT	14
3.1 Bytecode To Jimple	14
3.2 Compiling & Running Soot	14
3.3 Jimple Examples	15
3.3.1 Hello World	15
3.3.2 Inheritance Test	17
4. FOUR APPROACHES	19
4.1 One Thread Per Method	19
4.2 One Thread Per Class	20
4.3 Fork-Join Framework	21
4.4 Multiple Classes Per Thread	23
5. LOCKING	25
5.1 Doop Side	25
5.1.1 CSVDatabase	25
5.1.2 Repesantation	25
5.2 Soot Side	25
5.2.1 Type Assigner	26
5.2.2 Pack Manager	26
5.2.3 Shimple -ssa	26
6. TIME RESULTS	27
ΠΙΝΑΚΑΣ ΟΡΟΛΟΓΙΑΣ	28

ABBREVIATIONS	29
REFERENCES	30

ΚΑΤΑΛΟΓΟΣ ΣΧΗΜΑΤΩΝ

Σχήμα 1:	Compiling Soot	14
Σχήμα 2:	Generating Jimple from .class	15
Σχήμα 3:	Generating Jimple from .jar	15
Σχήμα 4:	HelloWorld.java	15
Σχήμα 5:	HelloWorld.jimple	16
Σχήμα 6:	inheritanceTest.java	17
Σχήμα 7:	inheritanceTest.jimple	18
Σχήμα 8:	Linear Fact Generation	19
Σχήμα 9:	One Thread Per Method	20
Σχήμα 10:	One Thread Per Class	21
Σχήμα 11:	Fork-Join Basic-Use	22
Σχήμα 12:	Fork-Join Framework	22
Σχήμα 13:	Multiple Classes Per Thread	24
Σχήμα 14:	CSVDatabase.java	25
Σχήμα 15:	Repesantation.java	25
Σχήμα 16:	JimpleBodyPack.java	26
Σχήμα 17:	PackManager.java	26
Σχήμα 18:	Fact Generation Time Results	27

ΚΑΤΑΛΟΓΟΣ ΠΙΝΑΚΩΝ

Πίνακας 1: Soot 2.5.0 times	13
---------------------------------------	----

PROLOGUE

Το παρόν έγγραφο δημιουργήθηκε στην Αθήνα, το 2016

blahblahblahblah

blahblahblahblah

1. INTRODUCTION

eisagwgh gia doop kai soot

2. DOOP

Doop is a framework for pointer, or points-to, analysis of Java programs. Doop implements a range of algorithms, including context insensitive, call-site sensitive, and object-sensitive analyses, all specified modularly as variations on a common code base. [1]

2.1 Fact Generation

Doop before running a pointer or points-to analysis, integrates with Soot to generate the facts. Facts are in Jimple (**J**ava **s**imple), a typed 3-address IR suitable for performing optimizations, it only has 15 statements. The main problem is that fact generation takes more than 50% of total execution time. Below are presented a few time examples of the sequential Doop Fact Generation and total execution time (both fact generation and analysis times).

2.2 Doop Time Examples

Soot 2.5.0	antlr.jar	hsqldb.jar	batik.jar
Fact Generation	1.16 min.	1.23 min.	2.26 min.
Total time	3.18 min.	3.21 min.	4.34 min.

Πίνακας 1: Soot 2.5.0 times

3. SOOT

Originally, Soot started off as a Java optimization framework. By now, researchers and practitioners from around the world use Soot to analyze, instrument, optimize and visualize Java and Android applications. In our case, we use Soot to generate Jimple (or Shimple) facts from Java bytecode to run a pointer or points-to analysis. [2]

3.1 Bytecode To Jimple

Soot is able to translate Java bytecode to a typed 3-address IR, Jimple. Jimple (**J**ava **s**imple) is a very convinient IR for performing optimizations, it only has 15 statements.

At first, soot translates bytecode to untyped Jimple and introduces new local variables. Then it inferres Types to the untyped jimple. The next step is to linearize all the expressions to statements that only refernce at most 3 local variables or constants.

In a general case the way Soot handles Java bytecode classes is the following: Soot is launched by specifying a directory with the Application code as a parameter. In this example only Java bytecode, either a class file or a jar. First, the *main()* method of the Main class is executed and calls *Scene.loadNecessaryClasses()*. This method loads basic Java classes and then loads specific Application classes by calling *loadClass()*. Then, *SootResolver.resolveClass()* is called. The resolver calls *SourceLocator.getClassSource()* to fetch a reference to a ClassSource, an interface between the file containing the Java bytecode and Soot. In our case the class source is a *CoffiClassSource* because it is the coffi module which handles the conversion from Java bytecode to Jimple. Then, the resolver having a reference to a class source, calls *resolve()* on it. This methods in turn calls *soot.coffi.Util.resolveFromClassFile()* which creates a *SootClass* from the corresponding Java bytecode class. All source fields of Soot class methods are set to refer to a *CoffiMethodSource* object. This object is used later to get the Jimple representation of the method. For example, if during an analysis with Soot the analysis code calls *SootMethod.getActiveBody()* and the Jimple code of the method was not already generated, *getActiveBody()* will call *CofficMethodSource.g* to compute Jimple code from the Java bytecode. The Jimple code representation of the method can then be analyzed and/or transformed.

3.2 Compiling & Running Soot

1	ant	/* To compile */
2	ant classesjar	/* To generate the sootclasses jar file */
3	ant fulljar	/* To generate the complete soot jar file */

Σχήμα 1: Compiling Soot

```
1      (create a test.java)
2      javac test.java
3      java -cp ./lib/soot-trunk.jar soot.Main -f J -cp
        ./usr/lib/jvm/java-7-openjdk-amd64/jre/lib/rt.jar test
```

Σχήμα 2: Generating Jimple from .class

```
1      java -cp soot-trunk.jar soot.Main -f J -cp
        ./usr/lib/jvm/java-7-openjdk-amd64/jre/lib/rt.jar -process-dir
        pathtotest.jar
```

Σχήμα 3: Generating Jimple from .jar

3.3 Jimple Examples

Below are two simple java programs along with their jimple translation. The first one is the classic HelloWorld, and the second is a simple inheritance test that depends on the user's input. The local variables which start with a \$ sign represent stack positions and not local variables in the original program whereas those without \$ represent real local variables.

3.3.1 Hello World

```
1      public class helloWorld {
2          public static void main(String[] args) {
3              System.out.println("Hello, World");
4          }
5      }
```

Σχήμα 4: HelloWorld.java

```
1      public class helloWorld extends java.lang.Object {
2
3          public void <init>() {
4              helloWorld r0;
5              r0 := @this: helloWorld;
6              specialinvoke r0.<java.lang.Object: void <init>()>();
7              return;
8          }
9
10         public static void main(java.lang.String []) {
11             java.lang.String [] r0;
12             java.io.PrintStream $r1;
13             r0 := @parameter0: java.lang.String[];
14             $r1 = <java.lang.System: java.io.PrintStream out>;
15             virtualinvoke $r1.<java.io.PrintStream: void
16                 println (java.lang.String)>("Hello, World");
17             return;
18         }
19     }
```

Σχήμα 5: HelloWorld.jimple

3.3.2 Inheritance Test

```
1      public class inheritanceTest {
2
3          public static void main(String[] args) {
4              testA a;
5              if (args.length < 1) {
6                  a = new testA(5);
7              } else {
8                  a = new testB(5);
9              }
10             int result = a.getA();
11             System.out.println("the value of a is " + result);
12         }
13
14         public static class testA {
15             int a;
16
17             public testA(int a) {
18                 this.a = a;
19             }
20
21             public int getA() {
22                 return this.a;
23             }
24         }
25
26         public static class testB extends testA {
27             public testB(int a) {
28                 super(a+100);
29             }
30         }
31     }
32
```

Σχήμα 6: inheritanceTest.java

```

1      public class inheritanceTest extends java.lang.Object {
2
3          public void <init>() {
4              inheritanceTest r0;
5              r0 := @this: inheritanceTest;
6              specialinvoke r0.<java.lang.Object: void <init>()>();
7              return;
8          }
9
10         public static void main(java.lang.String []) {
11             java.lang.String [] r0;
12             int $i0, i1;
13             inheritanceTest$testA $r1, r2;
14             inheritanceTest$testB $r3;
15             java.io .PrintStream $r4;
16             java.lang.StringBuilder $r5, $r6, $r7;
17             java.lang.String $r8;
18             r0 := @parameter0: java.lang.String[];
19             $i0 = lengthof r0;
20             if $i0 >= 1 goto label1;
21             $r1 = new inheritanceTest$testA;
22             specialinvoke $r1.<inheritanceTest$testA: void <init>(int)>(5);
23             r2 = $r1;
24             goto label2;
25         label1:
26             $r3 = new inheritanceTest$testB;
27             specialinvoke $r3.<inheritanceTest$testB: void <init>(int)>(5);
28             r2 = $r3;
29         label2:
30             i1 = virtualinvoke r2.<inheritanceTest$testA: int getA()>();
31             $r4 = <java.lang.System: java.io.PrintStream out>;
32             $r5 = new java.lang.StringBuilder;
33             specialinvoke $r5.<java.lang.StringBuilder: void <init>()>();
34             $r6 = virtualinvoke $r5.<java.lang.StringBuilder:
35                 java.lang.StringBuilder append(java.lang.String)>("the value
36                 of a is ");
37             $r7 = virtualinvoke $r6.<java.lang.StringBuilder:
38                 java.lang.StringBuilder append(int)>(i1);
39             $r8 = virtualinvoke $r7.<java.lang.StringBuilder: java.lang.String
40                 toString()>();
41             virtualinvoke $r4.<java.io.PrintStream: void
42                 println (java.lang.String)>($r8);
43             return;
44         }
45     }

```

Σχήμα 7: inheritanceTest.jimple

4. FOUR APPROACHES

We now describe the main idea of Fact Generation from Dooop side. Given all the classes (sootClasses) to generate, Dooop iterates each one of them; writes all the superClasses, if exist, and then generate all fields (sootFields) and methods (sootMethods).

```

1      public class FactGenerator {
2          /* ... */
3
4          public void generate(sootClass) {
5              if (c.hasSuperclass() && !c.isInterface())
6                  _writer.writeDirectSuperclass(c, c.getSuperclass());
7              for(SootField f : c.getFields())
8                  generate(f);
9              for(SootMethod m : c.getMethods()) {
10                 Session session = new Session();
11                 generate(m, session);
12             }
13         }
14
15         public void generate(SootMethod m, Session session) {
16             /* ... */
17
18             /* This instruction spends more than 80% of FG time */
19             m.retrieveActiveBody()
20
21             /* ... */
22         }
23
24         /* ... */
25     }

```

Σχήμα 8: Linear Fact Generation

Having the previous basic structure in mind, and considering that *m.retrieveActiveBody()* spends more than 80% of total fact generation time, we tried to parallelize the method which calls *m.retrieveActiveBody()*. In order to do that, we approached the problem in four ways. The three of them are pretty much similar while the other one is based on a recursive Java Framework (Fork/Join Framework).

4.1 One Thread Per Method

Our first approach to parallelize Fact Generation is similar as the sequential one, but instead of having a loop over all Soot Methods and call *generate(m, session)*, we assign the task to a thread for each one of them. We created a new Java class, MethodGenerator, which is identical to FactGenerator and also has a *run()* method to generate sootMethods.

```

1      public class FactGenerator {
2          private ExecutorService MgExecutor = new ThreadPoolExecutor(8, 16, 0L,
3              TimeUnit.MILLISECONDS, new LinkedBlockingQueue<Runnable>());
4              /* ... */
5
6          public void generate(sootClass) {
7              if (c.hasSuperclass() && !c.isInterface())
8                  _writer.writeDirectSuperclass(c, c.getSuperclass());
9              for (SootField f : c.getFields())
10                  generate(f);
11              for (SootMethod m : c.getMethods()) {
12                  Session session = new Session();
13                  Runnable mg = new MethodGenerator();
14                  MgExecutor.execute(mg);
15              }
16          }
17
18      public class MethodGenerator {
19          public void run() {
20              generate(this.m, this.s)
21          }
22
23          /* ... */
24      }

```

Σχήμα 9: One Thread Per Method

4.2 One Thread Per Class

In our second approach, we tried to find out ways to gain more speedup. So, we observed that some threads did not have much work to do and finishing their task instantly. Allocating a new object and assigning to it a task just to finish instantly was an overhead. As a result, we tried to feed the threads more than just a method, we created a new thread for each class, not for each method.

```

1      public class FactGenerator {
2          private ExecutorService CgExecutor = new ThreadPoolExecutor(8, 16, 0L,
3              TimeUnit.MILLISECONDS, new LinkedBlockingQueue<Runnable>());
4              /* ... */
5
6          public void generate(sootClass) {
7              Runnable cg = new ClassGenerator();
8              CgExecutor.execute(cg);
9          }
10     }
11
12     public class ClassGenerator {
13         public void run() {
14             if (c.hasSuperclass() && !c.isInterface())
15                 _writer.writeDirectSuperclass(c, c.getSuperclass());
16             for(SootField f : c.getFields())
17                 generate(f);
18             for(SootMethod m : c.getMethods()) {
19                 Session session = new Session();
20                 Runnable mg = new MethodGenerator();
21                 MgExecutor.execute(mg);
22                 generate(m, session);
23             }
24         }
25         /* ... */
26     }

```

Σχήμα 10: One Thread Per Class

4.3 Fork-Join Framework

From Oracle's Java Documentation:

"The fork/join framework is an implementation of the ExecutorService interface that helps you take advantage of multiple processors. It is designed for work that can be broken into smaller pieces recursively. The goal is to use all the available processing power to enhance the performance of your application.

The center of the fork/join framework is the ForkJoinPool class, an extension of the AbstractExecutorService class. ForkJoinPool implements the core work-stealing algorithm and can execute ForkJoinTask processes.

The idea of using the fork/join framework is to write code that performs a segment of the work. The basic structure should be like the following pseudocode." [3]

```

1      if (my portion of the work is small enough) {
2          do the work directly
3      } else {
4          split my work into two pieces
5          invoke the two pieces and wait for the results
6      }

```

Σχήμα 11: Fork-Join Basic-Use

```

1      public class FactGenerator {
2          private ForkJoinPool classGeneratorPool = new ForkJoinPool();
3          /* ... */
4          public void generate(sootClass) {
5              if (c.hasSuperclass() && !c.isInterface())
6                  _writer.writeDirectSuperclass(c, c.getSuperclass());
7              for(SootClass i : c.getInterfaces())
8                  _writer.writeDirectSuperinterface(c, i);
9              for(SootField f : c.getFields())
10                 generate(f);
11              if (c.getMethods().size() > 0) {
12                  ClassGenerator classGenerator = new ClassGenerator(_writer,
13                      _ssa, c, 0, c.getMethods().size());
14                  classGeneratorPool.invoke(classGenerator);
15              }
16          }
17      }
18
19      public class ClassGenerator {
20          /* ... */
21          public void compute() {
22              List<SootMethod> sootMethods = _sootClass.getMethods();
23              /* if (my portion of the work is small enough) */
24              if (_to - _from < threshold) {
25                  for (int i = _from ; i < _to ; i++) {
26                      SootMethod m = sootMethods.get(i);
27                      Session session = new Session();
28                      generate(m, session);
29                  }
30              } else { /* split work*/
31                  int half = (_to - _from)/2;
32                  ClassGenerator c1 = new ClassGenerator(_writer, _ssa,
33                      _sootClass, _from, _from + half);
34                  ClassGenerator c2 = new ClassGenerator(_writer, _ssa,
35                      _sootClass, _from + half, _to);
36                  invokeAll(c1, c2);
37              }
38          }
39          /* ... */
40      }

```

Σχήμα 12: Fork-Join Framework

4.4 Multiple Classes Per Thread

Our final approach is similar as the second one, but instead of having one thread per class, we now have one thread per multiple classes. Even in the second approach some threads did not have much work to do. Below is presented in an abstract way the final stage of the code implementing the *Multiple Classes Per Thread* approach.

Getting away from the overhead produced by assignments and allocations of the first and second approach, this one had so far the best time results.

```

1      public class Driver {
2          public Driver(ThreadFactory factory, boolean ssa, int totalClasses) {
3              _factory = factory;
4              _ssa = ssa;
5              _classCounter = 0;
6              _sootClasses = new ArrayList<>();
7              _totalClasses = totalClasses;
8              _cores = Runtime.getRuntime().availableProcessors();
9              _executor = new ThreadPoolExecutor(_cores/2, _cores, 0L,
                TimeUnit.MILLISECONDS, new
                LinkedBlockingQueue<Runnable>());
10         }
11
12         public void doInParallel( List<SootClass> sootClasses) {
13             for(SootClass c : sootClasses)
14                 generate(c);
15             _executor.shutdown();
16             _executor.awaitTermination(Long.MAX_VALUE,
                TimeUnit.NANOSECONDS);
17         }
18
19         void generate(SootClass _sootClass) {
20             _classCounter++;
21             _sootClasses.add(_sootClass);
22             if ((_classCounter % _classSplit == 0) || (_classCounter +
                _classSplit-1 >= _totalClasses)) {
23                 Runnable runnable = _factory.newRunnable(_sootClasses);
24                 _executor.execute(runnable);
25                 _sootClasses = new ArrayList<>();
26             }
27         }
28     }
29
30     public class ThreadFactory {
31         /* ... */
32         public Runnable newRunnable(List<SootClass> sootClasses) {
33             if (_makeClassGenerator)
34                 return new FactGenerator(_factWriter, _ssa, sootClasses);
35             else
36                 return new FactPrinter(_ssa, _toStdout, _outputDir, _printWriter ,
                sootClasses);
37         }
38     }

```

Σχήμα 13: Multiple Classes Per Thread

5. LOCKING

Along with threads come locks. An incorrect use of locking could have very negative(?) effects on the results. Having more locks than we actually needed would have led to a time result as slow as the sequential one. Having less, would have led to races and deadlocks. So, locks should be used very consciously.

Luckily, in our case we used very few locks in both Doop and Soot. Below we explain in more detail.

5.1 Doop Side

The only two things we had to lock to prevent races and deadlocks were an access to the output file and three methods that access a `sootMethod`.

5.1.1 CSVDatabase

Lock file to prevent more than one thread to write at the same time.

```

1      synchronized(predicateFile) {
2          Writer writer = getWriter(predicateFile);
3          addColumn(writer, arg, shouldTruncate);
4          for (Column col : args)
5              addColumn(writer.append(SEP), col, shouldTruncate);
6          writer . write (EOL);
7      }
```

Σχήμα 14: CSVDatabase.java

5.1.2 Representation

The other synchronization we had to provide was in three methods that were accessing `SootMethods` while threads were active.

```

1      public synchronized String signature(SootMethod) { /*...*/ }
2      public synchronized String handler(SootMethod, Trap, Session) { /*...*/ }
3      public synchronized String compactMethod(SootMethod) { /*...*/ }
```

Σχήμα 15: Representation.java

5.2 Soot Side

The way Soot is implemented, it has a class that encloses all global objects (*G.java*). As we mentioned before, soot has a class that contains all global objects. Soot does much more than just translating bytecode to jimple, so it has various phases and a lot of different options for transformations given.

5.2.1 Type Assigner

The class *JimpleBodyPack* applies the transformations corresponding to the given options. In our case, it applies a *"jb.tr"* which means *"jimple body transformation"*. From bytecode to jimple translation this is the only pack needed, so we lock before applying Type Assigner and unlock afterwards.

```

1      lock.lock();
2      PackManager.v().getTransform("jb.tr").apply(b);
3      lock.unlock();

```

Σχήμα 16: *JimpleBodyPack.java*

5.2.2 Pack Manager

The class that manages the Packs containing the various phases and their options is *PackManager*. Lock before retrieving Class Hierarchy Analysis, unlock afterwards.

```

1      lock.lock();
2      p.add(new Transform("cg.cha", CHATransformer.v()));
3      p.add(new Transform("cg.spark", SparkTransformer.v()));
4      p.add(new Transform("cg.paddle", PaddleHook.v()));
5      lock.unlock();


```

Σχήμα 17: *PackManager.java*

5.2.3 Shimple -ssa

We also use Soot to provide **Shimple** which is an **SSA** variation of Jimple. Shimple generation is very similar to jimple but with a few differences. blahblabh prepei na to dw kiallo auto The class that handles the translation from jimple to shimple is *Shimple.java*. So far we have synchronized all Shimple-Body-creation methods.

6. TIME RESULTS



Σχήμα 18: Fact Generation Time Results

ΠΙΝΑΚΑΣ ΟΡΟΛΟΓΙΑΣ

Ακολουθεί δείγμα πίνακα ορολογίας.

κλάση	class
εντολή	command
περιβάλλον	environment

ABBREVIATIONS

NP	Non-deterministic polynomial time
----	-----------------------------------

REFERENCES

- [1] "Doop: Framework for Java Pointer Analysis" [Online] Available: <http://doop.program-analysis.org/>
- [2] "Sable: Soot" [Online] Available: <https://sable.github.io/soot/>
- [3] "Oracle Java Fork/Join Framework" [Online] Available: <https://docs.oracle.com/javase/tutorial/essential/concurrency/forkjoin.html>